

```
In [1]: import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt

from sklearn.model_selection import train_test_split

from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import RandomForestClassifier
from sklearn.linear_model import LogisticRegression

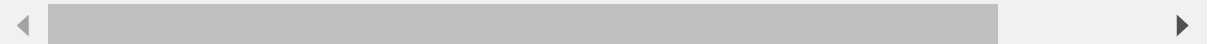
from sklearn.metrics import accuracy_score, classification_report
```

```
In [2]: data = pd.read_csv("C:/Users/User/Downloads/diabetes (2).csv")
```

```
In [3]: data.head()
```

Out[3]:

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	DiabetesPedigreeFunction
0	6	148	72	35	0	33.6	0.627
1	1	85	66	29	0	26.6	0.351
2	8	183	64	0	0	23.3	0.672
3	1	89	66	23	94	28.1	0.167
4	0	137	40	35	168	43.1	2.288



In [4]: data.describe

```
Out[4]: <bound method NDFrame.describe of
kinThickness  Insulin  BMI  \
0             6     148      72      35      0  33.6
1             1      85      66      29      0  26.6
2             8     183      64       0      0  23.3
3             1      89      66      23     94  28.1
4             0     137      40      35    168  43.1
..          ...     ...     ...     ...     ...     ...
763          10     101      76      48    180  32.9
764           2     122      70      27      0  36.8
765           5     121      72      23    112  26.2
766           1     126      60       0      0  30.1
767           1      93      70      31      0  30.4
```

```
DiabetesPedigreeFunction  Age  Outcome
0             0.627     50         1
1             0.351     31         0
2             0.672     32         1
3             0.167     21         0
4             2.288     33         1
..          ...     ...     ...
763          0.171     63         0
764          0.340     27         0
765          0.245     30         0
766          0.349     47         1
767          0.315     23         0
```

[768 rows x 9 columns]>

In [5]: data.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 768 entries, 0 to 767
Data columns (total 9 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   Pregnancies                          768 non-null   int64
1   Glucose                              768 non-null   int64
2   BloodPressure                        768 non-null   int64
3   SkinThickness                        768 non-null   int64
4   Insulin                              768 non-null   int64
5   BMI                                  768 non-null   float64
6   DiabetesPedigreeFunction             768 non-null   float64
7   Age                                  768 non-null   int64
8   Outcome                              768 non-null   int64
dtypes: float64(2), int64(7)
memory usage: 54.1 KB
```

In [6]: data.columns

```
Out[6]: Index(['Pregnancies', 'Glucose', 'BloodPressure', 'SkinThickness', 'Insulin',
              'BMI', 'DiabetesPedigreeFunction', 'Age', 'Outcome'],
              dtype='object')
```

```
In [7]: data.isnull()
```

```
Out[7]:
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	DiabetesPedigreeFunc	Outcome
0	False	False	False	False	False	False	False	0
1	False	False	False	False	False	False	False	0
2	False	False	False	False	False	False	False	0
3	False	False	False	False	False	False	False	0
4	False	False	False	False	False	False	False	0
...	...	...	...	...	...	...	...	...
763	False	False	False	False	False	False	False	0
764	False	False	False	False	False	False	False	0
765	False	False	False	False	False	False	False	0
766	False	False	False	False	False	False	False	0
767	False	False	False	False	False	False	False	0

768 rows × 9 columns



```
In [8]: data.nunique()
```

```
Out[8]: Pregnancies      17
Glucose      136
BloodPressure  47
SkinThickness  51
Insulin      186
BMI          248
DiabetesPedigreeFunction  517
Age          52
Outcome       2
dtype: int64
```

```
In [9]: np.sum(data.isnull().any(axis=1))
```

```
Out[9]: 0
```


```
In [10]: data.Pregnancies.value_counts()
```

```
Out[10]: Pregnancies
1      135
0      111
2      103
3       75
4       68
5       57
6       50
7       45
8       38
9       28
10      24
11      11
13      10
12       9
14       2
15       1
17       1
Name: count, dtype: int64
```

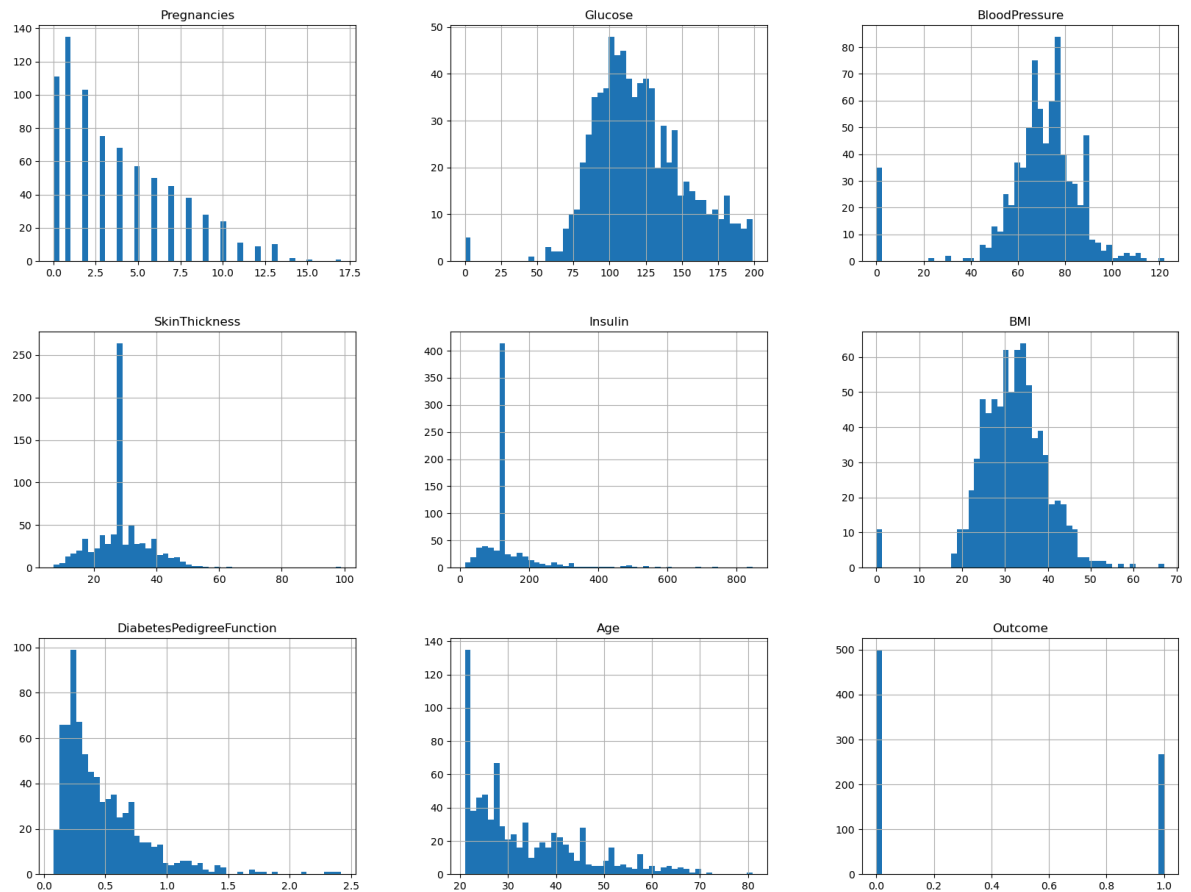
```
In [11]: data['SkinThickness'] = data['SkinThickness'].replace(0, np.nan) # Replace ze
data['Insulin'] = data['Insulin'].replace(0, np.nan) # Replace zero insulin l
data.fillna(data.median(), inplace=True)
```

```
In [12]: sns.countplot(x='Pregnancies', data=data)
```

```
Out[12]: <Axes: xlabel='Pregnancies', ylabel='count'>
```




```
In [13]: data.hist(bins =50, figsize=(20,15))
plt.show()
```



```
In [14]: corr_matrix = data.corr()
corr_matrix['Outcome'].sort_values(ascending=False)
```

```
Out[14]: Outcome      1.000000
Glucose      0.466581
BMI          0.292695
Age          0.238356
Pregnancies  0.221898
SkinThickness 0.214873
Insulin      0.203790
DiabetesPedigreeFunction 0.173844
BloodPressure 0.065068
Name: Outcome, dtype: float64
```

```
In [15]: plt.figure(figsize=(8,6))
sns.heatmap(corr_matrix, annot = True, cmap = 'coolwarm')
plt.show()
```



```
In [16]: y = data['Outcome']
```

```
In [17]: X = data.drop('Outcome', axis=1)
```

```
In [18]: #Splitting Data for training and testing in two parts
```

```
X_train,X_test,y_train,y_test = train_test_split(X,y,test_size =0.2, random_st
```

```
In [19]: from sklearn.preprocessing import StandardScaler
scaler =StandardScaler()
X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)
```

```
In [20]: dt_model = DecisionTreeClassifier()
dt_model.fit(X_train_scaled, y_train)
y_pred_dt = dt_model.predict(X_test_scaled)

# Evaluate Decision Tree
print("Decision Tree Accuracy:", accuracy_score(y_test, y_pred_dt))
print("Decision Tree Classification Report:\n", classification_report(y_test,
```

```
Decision Tree Accuracy: 0.6818181818181818
Decision Tree Classification Report:
              precision    recall  f1-score   support

      0       0.75       0.77       0.76       100
      1       0.55       0.52       0.53        54

 accuracy          0.68          154
 macro avg         0.65          154
 weighted avg      0.68          154
```

```
In [21]: rf_model = RandomForestClassifier()
rf_model.fit(X_train_scaled, y_train)
y_pred_dt = rf_model.predict(X_test_scaled)
# Evaluate Random Forest
print("Random Forest Accuracy:", accuracy_score(y_test, y_pred_dt))
print("Random Forest Classification Report:\n", classification_report(y_test,
```

```
Random Forest Accuracy: 0.7987012987012987
Random Forest Classification Report:
              precision    recall  f1-score   support

      0       0.84       0.85       0.85       100
      1       0.72       0.70       0.71        54

 accuracy          0.80          154
 macro avg         0.78          154
 weighted avg      0.80          154
```

```
In [22]: lr_model = LogisticRegression()
lr_model.fit(X_train_scaled, y_train)
y_pred_dt = lr_model.predict(X_test_scaled)
# Evaluate Logistic Regression
print("Logistic Regression Accuracy:", accuracy_score(y_test, y_pred_dt))
print("Logistic Regression Classification Report:\n", classification_report(y_
```

Logistic Regression Accuracy: 0.7922077922077922

Logistic Regression Classification Report:

	precision	recall	f1-score	support
0	0.82	0.87	0.84	100
1	0.73	0.65	0.69	54
accuracy			0.79	154
macro avg	0.77	0.76	0.77	154
weighted avg	0.79	0.79	0.79	154

In [ ]:

```
In [23]: data.to_csv('data.csv')
```